

Importing the library

```
In [1]:  
  
import numpy as np  
import pandas as pd  
import seaborn as sns  
import matplotlib.pyplot as plt  
  
import warnings  
warnings.filterwarnings('ignore')
```

Importing the dataset

```
In [2]:  
  
df=pd.read_csv(r"C:\Users\91956\Desktop\IRIS.csv")  
df
```

Out[2]:

	sepal_length	sepal_width	petal_length	petal_width	species
0	5.1	3.5	1.4	0.2	Iris-setosa
1	4.9	3.0	1.4	0.2	Iris-setosa
2	4.7	3.2	1.3	0.2	Iris-setosa
3	4.6	3.1	1.5	0.2	Iris-setosa
4	5.0	3.6	1.4	0.2	Iris-setosa
...
145	6.7	3.0	5.2	2.3	Iris-virginica
146	6.3	2.5	5.0	1.9	Iris-virginica
147	6.5	3.0	5.2	2.0	Iris-virginica
148	6.2	3.4	5.4	2.3	Iris-virginica
149	5.9	3.0	5.1	1.8	Iris-virginica

150 rows x 5 columns

EDA

Shape

In [3]:

```
df.shape
```

Out[3]:

```
(150, 5)
```

Columns and their types

In [4]:

```
df.columns
```

Out[4]:

```
Index(['sepal_length', 'sepal_width', 'petal_length', 'petal_width',  
      'species'],  
      dtype='object')
```

In [5]:

```
df.dtypes
```

Out[5]:

```
sepal_length    float64  
sepal_width     float64  
petal_length    float64  
petal_width     float64  
species         object  
dtype: object
```

In [6]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 150 entries, 0 to 149  
Data columns (total 5 columns):  
 #   Column          Non-Null Count  Dtype  
---  ---  
 0   sepal_length    150 non-null   float64
```

```
1  sepal_width  150 non-null  float64
2  petal_length 150 non-null  float64
3  petal_width  150 non-null  float64
4  species      150 non-null  object
dtypes: float64(4), object(1)
memory usage: 6.0+ KB
```

Null value analysis

In [7]:

```
df.isna().sum()
```

Out[7]:

```
sepal_length    0
sepal_width      0
petal_length     0
petal_width      0
species          0
dtype: int64
```

Duplicate values

In [8]:

```
df.duplicated().sum()
```

Out[8]:

3

In [9]:

```
df[df.duplicated()]
```

Out[9]:

	sepal_length	sepal_width	petal_length	petal_width	species
34	4.9	3.1	1.5	0.1	Iris-setosa
37	4.9	3.1	1.5	0.1	Iris-setosa
142	5.8	2.7	5.1	1.9	Iris-virginica


In [10]:

```
df.drop_duplicates(inplace=True)
```

Outlier's detection


In [11]:


```
sns.boxplot(y='sepal_length', data=df)  
plt.show()
```



In [12]:


```
sns.boxplot(y='sepal_width', data=df)  
plt.show()
```






In [13]:


```
sns.boxplot(y='petal_length', data=df)  
plt.show()
```



In [14]:

```
sns.boxplot(y='petal_width', data=df)  
plt.show()
```





Statistical summary

In [15]:

```
df.describe()
```


Out[15]:

	sepal_length	sepal_width	petal_length	petal_width
count	147.000000	147.000000	147.000000	147.000000
mean	5.856463	3.055782	3.780272	1.208844
std	0.829100	0.437009	1.759111	0.757874
min	4.300000	2.000000	1.000000	0.100000
25%	5.100000	2.800000	1.600000	0.300000
50%	5.800000	3.000000	4.400000	1.300000
75%	6.400000	3.300000	5.100000	1.800000
max	7.900000	4.400000	6.900000	2.500000

Correlation

In [23]:

```
sns.heatmap(df.corr(), cmap = "YlGnBu", annot = True)
plt.show()
```



Interpretation

In [16]:

```
df
```

Out[16]:


	sepal_length	sepal_width	petal_length	petal_width	species
0	5.1	3.5	1.4	0.2	Iris-setosa
1	4.9	3.0	1.4	0.2	Iris-setosa
2	4.7	3.2	1.3	0.2	Iris-setosa
3	4.6	3.1	1.5	0.2	Iris-setosa
4	5.0	3.6	1.4	0.2	Iris-setosa

4	5.0	3.6	1.4	0.2	Iris-setosa
sepal_length	sepal_width	petal_length	petal_width		species
...
145	6.7	3.0	5.2	2.3	Iris-virginica
146	6.3	2.5	5.0	1.9	Iris-virginica
147	6.5	3.0	5.2	2.0	Iris-virginica
148	6.2	3.4	5.4	2.3	Iris-virginica
149	5.9	3.0	5.1	1.8	Iris-virginica

147 rows x 5 columns


In [18]:

```
# Highest number of species according to petal_width
sns.barplot(x=df['species'],y=df['petal_width'])
plt.xticks(rotation=90)
plt.show()
```




In [19]:


```
# # Highest number of species according to petal_length
sns.barplot(x=df['species'], y=df['petal_length'])
plt.xticks(rotation=90)
plt.show()
```



In [20]:


```
#Highest number of species according to sepal_width
sns.barplot(x=df['species'], y=df['sepal_width'])
plt.xticks(rotation=90)
plt.show()
```

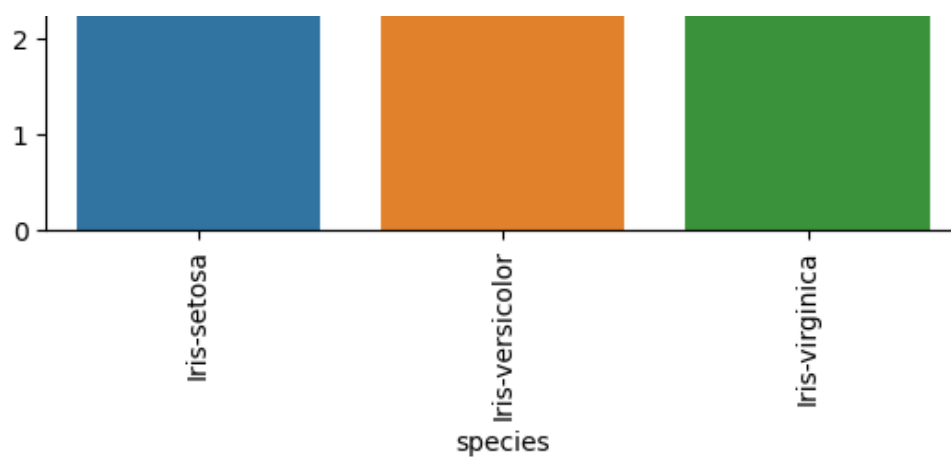




In [21]:

```
#Highest number of species according to sepal_length  
sns.barplot(x=df['species'], y=df['sepal_length'])  
plt.xticks(rotation=90)  
plt.show()
```





Replacing the values of object column with numbers

In [25]:

```
df['species']=df['species'].replace({'Iris-setosa':0,'Iris-versicolor':1,'Iris-virginica':2})
```

In [24]:

```
df.species.value_counts()
```

Out[24]:

```
Iris-versicolor    50
Iris-virginica     49
Iris-setosa        48
Name: species, dtype: int64
```

Machine learning

In [26]:

```
x=df.drop('species',axis=1)
y=df['species']
```

Splitting the dataset into training and testing

In [16]:

```
from sklearn.model_selection import train_test_split
```

```
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2,random_state=7)
```

Importing the algorithm and the performance measure

In [17]:

```
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, accuracy_score, confusion_matrix
```

In [18]:

```
classifier=LogisticRegression()  
lr=classifier.fit(x_train,y_train)  
y_pred=classifier.predict(x_test)  
y_pred_train=classifier.predict(x_train)
```

In [19]:

```
print(classification_report(y_test,y_pred))
print(classification_report(y_train,y_pred_train))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	7
1	0.83	0.83	0.83	12
2	0.82	0.82	0.82	11
accuracy			0.87	30
macro avg	0.88	0.88	0.88	30
weighted avg	0.87	0.87	0.87	30
	precision	recall	f1-score	support
0	1.00	1.00	1.00	43
1	1.00	0.97	0.99	38
2	0.97	1.00	0.99	39
accuracy			0.99	120
macro avg	0.99	0.99	0.99	120
weighted avg	0.99	0.99	0.99	120

In [20]:

[illegible]

```
dt=classifier2.fit(x_train,y_train)
y_pred=classifier2.predict(x_test)
y_pred_train=classifier2.predict(x_train)

print(classification_report(y_test,y_pred))
print(classification_report(y_train,y_pred_train))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	7
1	0.91	0.83	0.87	12
2	0.83	0.91	0.87	11
accuracy			0.90	30
macro avg	0.91	0.91	0.91	30
weighted avg	0.90	0.90	0.90	30

	precision	recall	f1-score	support
0	1.00	1.00	1.00	43
1	1.00	1.00	1.00	38
2	1.00	1.00	1.00	39
accuracy			1.00	120
macro avg	1.00	1.00	1.00	120
weighted avg	1.00	1.00	1.00	120

In [21]:

```
from sklearn.ensemble import RandomForestClassifier
classifier4=RandomForestClassifier(random_state=7, max_depth=7,
                                   criterion='gini',max_leaf_nodes=15,
                                   min_samples_split=40,n_estimators=30)

rf=classifier4.fit(x_train,y_train)
y_pred=classifier4.predict(x_test)
y_pred_train=classifier4.predict(x_train)

print(classification_report(y_test,y_pred))
print(classification_report(y_train,y_pred_train))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	7
1	0.83	0.83	0.83	12
2	0.82	0.82	0.82	11
accuracy			0.87	30
macro avg	0.88	0.88	0.88	30
weighted avg	0.87	0.87	0.87	30

	precision	recall	f1-score	support
0	1.00	1.00	1.00	43
1	0.95	1.00	0.97	38
2	1.00	0.95	0.97	39
accuracy			0.98	120
macro avg	0.98	0.98	0.98	120
weighted avg	0.98	0.98	0.98	120

In [22]:

```
from sklearn.svm import SVC
svc=SVC(random_state=7)
svm=svc.fit(x_train,y_train)
y_pred=svc.predict(x_test)
y_pred_train=svc.predict(x_train)

print(classification_report(y_test,y_pred))
print(classification_report(y_train,y_pred_train))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	7
1	0.83	0.83	0.83	12
2	0.82	0.82	0.82	11
accuracy			0.87	30
macro avg	0.88	0.88	0.88	30
weighted avg	0.87	0.87	0.87	30

	precision	recall	f1-score	support
0	1.00	1.00	1.00	43
1	1.00	0.97	0.99	38
2	0.97	1.00	0.99	39
accuracy			0.99	120
macro avg	0.99	0.99	0.99	120
weighted avg	0.99	0.99	0.99	120

In [23]:

```
from sklearn.ensemble import AdaBoostClassifier
adbc=AdaBoostClassifier(random_state=7)
adbcl=adbc.fit(x_train,y_train)
y_pred=adbc.predict(x_test)
```

```
y_pred_train=adbc.predict(x_train)

print(classification_report(y_test,y_pred))
print(classification_report(y_train,y_pred_train))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	7
1	0.83	0.83	0.83	12
2	0.82	0.82	0.82	11
accuracy			0.87	30
macro avg	0.88	0.88	0.88	30
weighted avg	0.87	0.87	0.87	30

	precision	recall	f1-score	support
0	1.00	1.00	1.00	43
1	0.95	1.00	0.97	38
2	1.00	0.95	0.97	39
accuracy			0.98	120
macro avg	0.98	0.98	0.98	120
weighted avg	0.98	0.98	0.98	120